



BeatAML imputation experiment — mean per-mutation ACCURACY (not AUROC)



Accuracy is prevalence-dominated (mean prevalence 5.6%): the trivial call-all-absent classifier scores 0.944, beating every model at its operating point — so accuracy barely discriminates here. Still, the ordering matches AUROC: imputed $\approx$ random ($\Delta+0.002$, $p=0.57$) $\rightarrow$ nonlinear feature-expansion, not biology.