






Per-mutation precision vs recall, by model and dataset



Each dot = one mutation

- n⁺ = 5 positives
- n⁺ = 20 positives
- n⁺ = 50 positives

Top-right = ideal. Dot size = positives available. Dashed = panel means. Precision is the discriminating axis for rare mutations (few positives, many chances to false-positive).

Multimodal at its nested-CV (honest) threshold; bulk model at its deployed F1-max threshold. Precision falls with prevalence, so rare drivers sit low even when their ranking (AUROC) is strong.